

# CUDA Streams, Events and asynchronous memory copies

*Jochen Kreutz*

**GPU Programming@Jülich Supercomputing Centre | Jülich | 25-27 April 2016**


## Overview

- Pinned (pagelocked) host memory
- Asynchronous and concurrent memory copies
- CUDA streams
  - The default stream and the `cudaStreamNonBlocking` flag
- CUDA Events
- CUBLAS
- nvprof + nvvp recap

## Pinned Host Memory

- Host memory allocated with `malloc` is pagable
  - Memory pages associated with the memory can be moved around by the OS Kernel, e.g. to swap space on hard disk
- Transfers to and from the GPU memory need to go over PCI-E
  - PCI-E transfers are handled by DMA engines on the GPU and work independently of the CPU/OS kernel
    - If OS kernel moves memory pages involved in such a DMA transfer the wrong data will be moved
    - Pinning memory pages inhibit the OS kernel from moving them around and make them usable to DMA transfer


# Pinned Host Memory



Host memory allocated with **malloc** is staged through a pinned memory buffer managed by the CUDA Driver

- No asynchronous memory copies are possible (CPU interaction is necessary to drive the pipeline)
- Higher latency and lower bandwidth compared to DMA transfers

## Pinned Host Memory



### Using pinned host memory

- enables asynchronous memory copies
- Lowers latency and increases bandwidth

## Pinned Host Memory – How to use it?

- Using POSIX functions like **mlock** is not sufficient, because the CUDA driver needs to know that the memory is pinned
- Two ways to get pinned host memory
  - Using **cudaMallocHost/cudaFreeHost** to allocate new memory
  - Using **cudaHostRegister/cudaHostUnregister** to pin memory after allocation
- **cudaMemcpy** makes automatic use of it
- **cudaMemcpyAsync** can be used to issue asynchronous memory copies
- Can be directly accessed from Kernels (zero-copy) – use **cudaHostGetDevicePointer**

## CUDA Streams

- CUDA Streams are work queues to express concurrency between different tasks, e.g.
  - host to device memory copies
  - device to host memory copies
  - kernel execution
- To overlap different tasks just launch them in different streams
  - All tasks launched into the same stream are executed in order
  - Tasks launched into different streams might execute concurrently (depending on available resources: two copy engines, compute resources)

# CUDA Streams – How to use them?

## ■ Create/Destroy

```
cudaStream_t stream;  
cudaStreamCreate( &stream );  
cudaStreamDestroy(stream);
```

## ■ Launch

```
my_kernel<<<grid,block,0,stream>>>(...);  
cudaMemcpyAsync( ..., stream );
```

## ■ Synchronize

```
cudaStreamSynchronize( stream );
```




# CUDA Streams – The default (null) stream

- Kernel launches are always asynchronous

- Which stream is used here?

```
my_kernel<<<grid,block>>> (...);
```


- The default (null) stream



# CUDA Streams – The `cudaStreamNonBlocking`

## Watch for false dependencies!

- The default stream waits for work in all other streams which
  - Do not have the `cudaStreamNonBlocking` flag set



User streams with the `cudaStreamNonBlocking` flag set can execute concurrently to stream 0

## CUDA Events

- CUDA Events are synchronization markers that can be used to:
  - Time asynchronous tasks in streams
  - Allow fine grained synchronization within a stream
  - Allow inter stream synchronization, e.g. let a stream wait for an event in another stream

# CUDA Events – How to use them?

## ■ Create/Destroy

```
cudaEventCreate( &event );  
cudaEventDestroy( event );
```

## ■ Record

```
cudaEventRecord( event, stream );
```

## ■ Query

```
cudaEventQuery( event );
```

## ■ Synchronize

```
cudaEventSynchronize( event );
```

## ■ Timing

```
cudaEventElapsedTime( &time, start, end );
```

# CUDA Events – Example for Kernel Timing

## Kernel timing

```
cudaEventRecord ( startEvent, stream );  
my_kernel<<<grid,block,0,stream>>>(...);  
cudaEventRecord ( endEvent, stream );  
  
//Host can do other work  
  
//Get runtime of my_kernel in ms  
float runtime = 0.0f;  
cudaEventSynchronize ( endEvent );  
cudaEventElapsedTime ( &runtime, startEvent, endEvent );
```

# Calling CUBLAS

## How to use CUBLAS function

```
#include "cublas_v2.h"

...

cublasHandle_t handle;

//Initialize cuBLAS
cublasCreate(&handle);

//Set cuBLAS execution stream
cublasSetStream(handle, stream);

//Call SAXPY
cublasSaxpy(handle, n, &alpha, x, 1, y, 1);
...
//Free up resources
cublasDestroy(handle);
```

## The command line profiler nvprof

- Simple launcher to get profiles of your application
- Profiles CUDA Kernels and API calls

### How to use command line profiler

```
> nvprof ./jacobi


===== NVPROF is profiling jacobi...
===== Command: jacobi
Jacobi (serial)
[...] snip
===== Profiling result:
Time(%)      Time      Calls      Avg      Min      Max      Name
 72.14    352.65ms    1000    352.65us  350.48us  354.94us  Jacobi_86_gpu
 26.02    127.23ms    1000    127.23us   93.48us  128.34us  Jacobi_74_gpu
  0.84      4.09ms    1000      4.09us   4.04us   4.36us  Jacobi_96_gpu_red
  0.61      3.00ms    1009      2.97us   2.78us  56.16us  [CUDA memcpy HtoD]
  0.39      1.91ms    1002      1.91us   1.82us  52.41us  [CUDA memcpy DtoH]
```

# nvprof interoperability with nvvp

- nvprof can write the application profile to nvvp compatible file:

```
nvprof -o jacobi.nvprof ./jacobi
```

- Import in nvvp





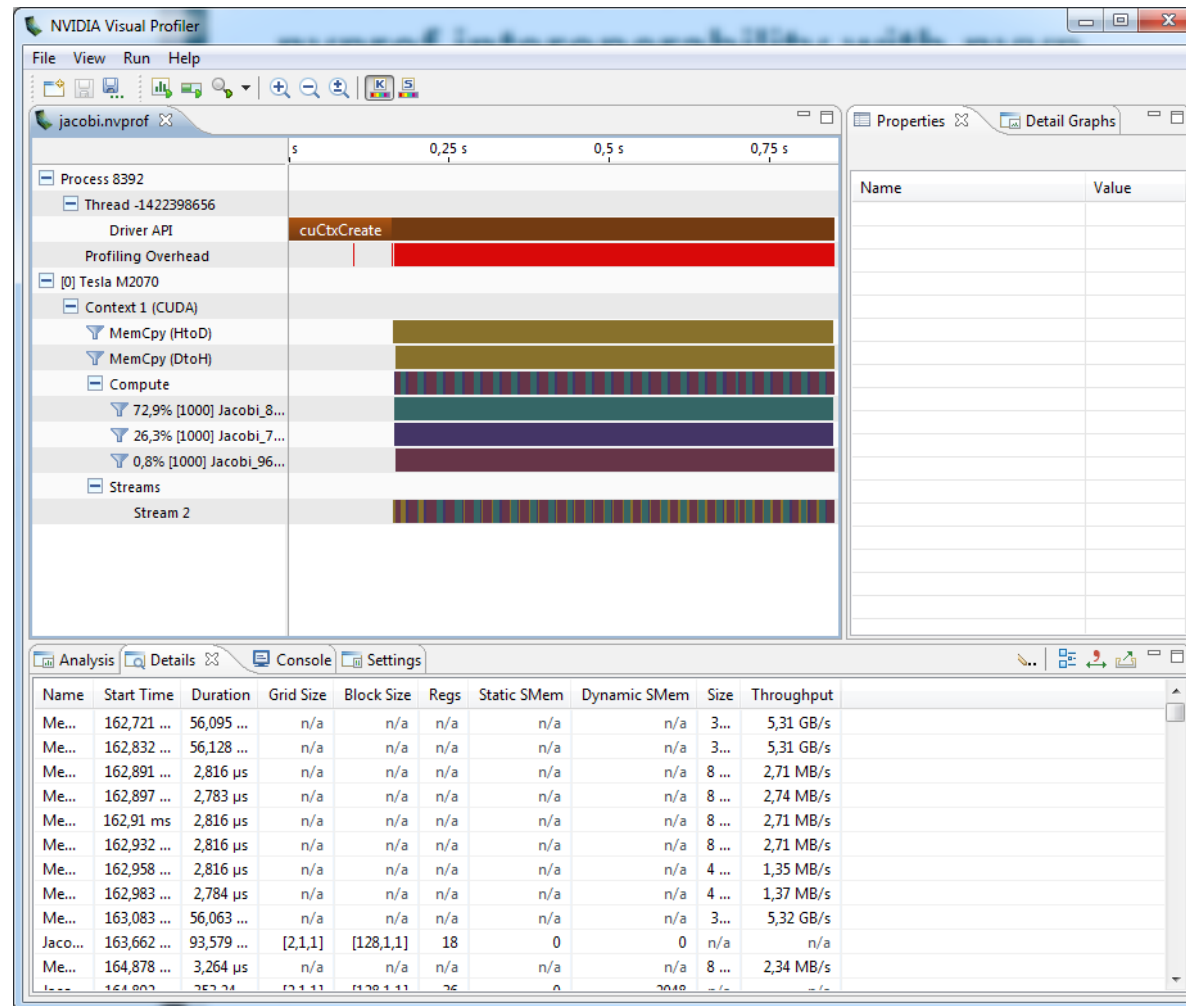
# nvprof important command-line options

## How to use command line profiler

### Options:

- `-o, --output-profile <filename>`  
Output the result file which can be imported later or opened by the NVIDIA Visual Profiler.
- `--events <event names>`  
Specify the events to be profiled on certain device(s). Multiple event names separated by comma can be specified. Which device(s) are profiled is controlled by the '--devices' option. Otherwise events will be collected on all devices. For a list of available events, use '--query-events'.
- `--query-events`  
List all the events available on each device.
- `-h, --help`  
Print this help information.

# nvvp



# Task 1: Optimize and overlap host to device and device to host transfers

## ■ Task 1a:

- Follow TODOs in *CUDAStreams/exercises/tasks/task1a.cu*
  - Allocate host buffers in pinned memory
- View nvprof profile in nvvp

## ■ Task 1b:

- Follow TODOs in *CUDAStreams/exercises/tasks/task1b.cu*
  - Create Upload and Download Stream
  - Issue Host to Device and Device to Host Transfer asynchronously in the two new streams.
- View nvprof profile in nvvp

## ■ Solutions in *CUDAStreams/exercises/solutions*

## ■ Slides are in

*CUDAStreams/slides/CUDAStreams\_and\_Events.pdf*

## Task 2:

- Follow TODOs in *CUDAStreams/exercises/tasks/task2.cu*
  - Set CUBLAS execution stream
  - Call CUBLAS SAXPY
  - Fix position of `cudaStreamSynchronize`
- Solution in *CUDAStreams/exercises/solutions/task2.cu*
- Slides are in  
*CUDAStreams/slides/CUDAStreams\_and\_Events.pdf*